

MIRA — 3-Minute Winning Demo Script (ACU x AWS Hackathon)

Team: Vivek · Narendra · Diwash **Scenario we demo:** an internationally-trained IT professional practising a job interview with a skeptical hiring manager — a struggle the three of us know first-hand. **Hard rules:** finish in **3:00**. Every sentence earns a rubric point. Never wait on a slow model reply on camera — keep narrating, cut to the report (that's the payoff).

Rubric coverage: Functionality /30 · UX /20 · Innovation /20 · Educational Impact /20 · Teamwork /10.

Speaking order & roles

- **Vivek** — opening (what we built + security-first design) + close. (the narrative)
 - **Narendra** — drives the app: types the scenario, speaks the interview turn. (the live demo)
 - **Diwash** — walks the coaching report + downloads the recording. (the payoff)
-

THE SCRIPT (with hooks →)

0:00–0:35 — VIVEK — WHAT WE BUILT + SECURITY-FIRST DESIGN → INNOVATION, FUNCTIONALITY

"This is **MIRA** — a virtual client-interaction web app, and we built it **privacy-first**. The entire conversation runs **on this machine**: the dialogue comes from a **local AI model (Llama 3.1)**, the client's voice is a **local speech model (Kokoro)**, and speech-to-text runs **in the browser**. **Nothing the student says ever leaves the device** — that's the safety guarantee that makes this usable for real, sensitive practice. For the one thing the cloud does best — computer vision — we use **AWS Amazon Rekognition** to analyse the student's eye contact.

In one app we built: **scenario-driven 3D environments**, a **live, fully-local spoken conversation**, **on-screen captions**, **eye-contact and communication analytics**, and a **downloadable recording** of the whole session.

Let me show you with a real scenario every international student knows — landing that first IT job. Narendra."

0:30–1:05 — NARENDRA — SCENARIO DRIVES THE ENVIRONMENT → *FUNCTIONALITY* /30

- Type into the box:

"I'm an internationally-trained IT professional with years of experience back home, but I keep getting rejected here. I want to practise a job interview with a hiring manager who is skeptical about my overseas experience."

- Click **Begin**.

"From that one sentence MIRA builds the whole scene — a professional office, a skeptical hiring-manager persona with a hidden concern, mood-driven lighting, and a matching voice. And there I am, on the screen in the room — the camera puts the student and the client face to face, exactly as the brief asks."

1:05–1:55 — NARENDRA — LIVE INTERACTION → *FUNCTIONALITY, INNOVATION, UX* /20

- Push to talk, say one strong line:

"Thanks for seeing me. I know my experience is from overseas, but I led a team of four building production systems — let me show you how that transfers."

- The client replies **aloud**. Point at the screen:

"Everything you just heard is **local**: the reply came from a **Llama 3.1 model on this machine**, and the voice is **Kokoro, a local speech model** — no audio, no transcript, nothing about my answer is sent to any cloud or AI company. You can read it in the **live captions**, and watch the lighting warm up as I build rapport — the client's **mood** is shifting. The only thing touching the cloud is **AWS Amazon Rekognition**, checking my eye contact through the webcam."

1:55–2:35 — DIWASH — THE COACHING REPORT → *EDUCATIONAL IMPACT* /20

- Click **End session & get report**. Walk it fast:

"Here's what makes this a learning tool, not a toy. An instant coaching report: an overall score and skills radar; **measured signals** — talk-ratio, filler words, pace, and **eye-contact percentage from Amazon Rekognition**; and critically — **every score is backed by a direct quote from what I actually said**. Plus two specific tips for next time. That's the reflective-practice loop: act, see real evidence, improve."

2:35–3:00 — DIWASH → VIVEK — RECORDING + CLOSE → FUNCTIONALITY, TEAMWORK /10

- Diwash clicks **Download recording**, show the `.webm` plays back the scene **with captions and both voices**.

(Diwash) “Full audio and video of the session — captions burned in — downloadable in one click, up to 30 minutes.”

- **Vivek closes:**

”So: a scenario-driven 3D environment, a **fully on-device conversation** — local dialogue, local voice, in-browser speech, so sensitive practice data never leaves the machine — with **AWS Rekognition** adding objective eye-contact analytics, and an evidence-based coaching report that actually helps you improve. That’s MIRA — practice that’s **safe, private, and real**. Thank you.”



How to produce the video (do this now)

Tooling: record the screen with the running app at `localhost:3000`.

- Mac: **Cmd+Shift+5** → “Record Selected Portion” over the browser window (records screen; use the mic for narration). Or **OBS** if you want a webcam talking-head in a corner.
- Capture system audio too so the client’s voice is in the video (OBS, or QuickTime “New Screen Recording” with audio). If unsure, narrate over it — the captions carry the meaning.

Process (≈15 min):

1. **Warm up** the model: run one throwaway turn first so the on-camera reply is fast.
2. **Dry run once** end-to-end with a stopwatch — trim until you land under 3:00.
3. **Record the final take.** One continuous screen-record is fine; one person narrates and clicks, the other two record their lines (intro/close) — or do it live together on a Teams call and record that.
4. Keep the **report screen on-screen longest** — it’s the highest-scoring moment.
5. Export as `MIRA-demo.mp4`, watch it once, confirm audio is audible and it’s < 3:00.

Shot checklist: scenario typed → office + avatar + your camera-in-scene → one spoken turn + voiced reply + captions + mood shift → report (eye-contact %, quotes, tips) → download recording.

Judge Q&A prep (1 minute — own these)

- **"Did you use AWS?"** → "Yes — **Amazon Rekognition** for eye-contact analysis, live. We made a deliberate **security** decision to run the whole conversation on-device — dialogue (local Llama), voice (local Kokoro), and speech-to-text (in-browser) — so sensitive practice data never leaves the machine. We use the cloud for vision, where it adds the most value."
 - **"Why run the AI locally — isn't cloud better?"** → "For healthcare, counselling and career coaching, the conversation is sensitive. Keeping the dialogue *and the voice* on-device means a university can actually deploy this under its data-governance rules. **Privacy and security are the feature** — and it also means zero per-conversation cloud cost and it works offline."
 - **"Is the scenario really driving the environment?"** → "The scenario text generates the persona live — role, hidden concern, mood, voice — and that drives the 3D room, the avatar, the lighting, and the Polly voice. Different scenario, different room and person."
 - **"How are the scores trustworthy?"** → "Numeric metrics are deterministic pure functions of the transcript — reproducible — and every qualitative score cites an exact quote. Eye contact is real Rekognition face-pose data."
 - **"What's the stack?"** → "Next.js + TypeScript + React Three Fiber; a service layer with Mock / Local / AWS implementations, so it runs offline on localhost and uses live AWS for voice and vision."
-

Pre-record checklist

- ☐ `npm run dev` fresh; **hard-refresh** (Cmd+Shift+R).
- ☐ Camera + mic allowed; good lighting; face the screen for a real eye-contact number.
- ☐ **Recording verified** — confirm the downloaded `.webm` has picture, captions, and BOTH voices BEFORE the final take.
- ☐ Model warmed (one throwaway turn); quiet room; other tabs closed.
- ☐ Final video is `MIRA-demo.mp4`, **audible, under 3:00**.